

ML1

July 18, 2024

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df=pd.read_csv('uber.csv')
```

```
[3]: df.head()
```

```
[3]:      Unnamed: 0      key  fare_amount  \
0    24238194    2015-05-07 19:52:06.0000003      7.5
1    27835199    2009-07-17 20:04:56.0000002      7.7
2    44984355    2009-08-24 21:45:00.00000061    12.9
3    25894730    2009-06-26 08:22:21.0000001      5.3
4    17610152    2014-08-28 17:47:00.000000188    16.0

      pickup_datetime  pickup_longitude  pickup_latitude  \
0  2015-05-07 19:52:06 UTC      -73.999817      40.738354
1  2009-07-17 20:04:56 UTC      -73.994355      40.728225
2  2009-08-24 21:45:00 UTC      -74.005043      40.740770
3  2009-06-26 08:22:21 UTC      -73.976124      40.790844
4  2014-08-28 17:47:00 UTC      -73.925023      40.744085

      dropoff_longitude  dropoff_latitude  passenger_count
0      -73.999512      40.723217      1
1      -73.994710      40.750325      1
2      -73.962565      40.772647      1
3      -73.965316      40.803349      3
4      -73.973082      40.761247      5
```

```
[4]: df.describe()
```

```
[4]:      Unnamed: 0      fare_amount  pickup_longitude  pickup_latitude  \
count  2.000000e+05  200000.000000    200000.000000    200000.000000
mean    2.771250e+07     11.359955      -72.527638      39.935885
```

std	1.601382e+07	9.901776	11.437787	7.720539
min	1.000000e+00	-52.000000	-1340.648410	-74.015515
25%	1.382535e+07	6.000000	-73.992065	40.734796
50%	2.774550e+07	8.500000	-73.981823	40.752592
75%	4.155530e+07	12.500000	-73.967154	40.767158
max	5.542357e+07	499.000000	57.418457	1644.421482

	dropoff_longitude	dropoff_latitude	passenger_count
count	199999.000000	199999.000000	200000.000000
mean	-72.525292	39.923890	1.684535
std	13.117408	6.794829	1.385997
min	-3356.666300	-881.985513	0.000000
25%	-73.991407	40.733823	1.000000
50%	-73.980093	40.753042	1.000000
75%	-73.963658	40.768001	2.000000
max	1153.572603	872.697628	208.000000

```
[5]: df.shape
```

```
[5]: (200000, 9)
```

```
[6]: df.dtypes
```

```
[6]: Unnamed: 0          int64
     key             object
     fare_amount      float64
     pickup_datetime  object
     pickup_longitude float64
     pickup_latitude  float64
     dropoff_longitude float64
     dropoff_latitude float64
     passenger_count   int64
     dtype: object
```

```
[7]: df=df.drop(['Unnamed: 0', 'key'], axis=1)
```

```
[8]: df.columns
```

```
[8]: Index(['fare_amount', 'pickup_datetime', 'pickup_longitude', 'pickup_latitude',
         'dropoff_longitude', 'dropoff_latitude', 'passenger_count'],
         dtype='object')
```

```
[9]: df.pickup_datetime=pd.to_datetime(df.pickup_datetime)
```

```
[10]: df.dtypes
```

```
[10]: fare_amount          float64
     pickup_datetime      datetime64[ns, UTC]
```

```

pickup_longitude      float64
pickup_latitude       float64
dropoff_longitude     float64
dropoff_latitude      float64
passenger_count       int64
dtype: object

```

```
[11]: df.isnull().sum()
```

```

[11]: fare_amount      0
      pickup_datetime  0
      pickup_longitude  0
      pickup_latitude  0
      dropoff_longitude 1
      dropoff_latitude 1
      passenger_count   0
      dtype: int64

```

```

[12]: df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].mean(),inplace=True)
      df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(),inplace=True)

```

```
[13]: df.isnull().sum()
```

```

[13]: fare_amount      0
      pickup_datetime  0
      pickup_longitude  0
      pickup_latitude  0
      dropoff_longitude 0
      dropoff_latitude  0
      passenger_count   0
      dtype: int64

```

```

[14]: df = df.assign(hour = df.pickup_datetime.dt.hour,
                    day = df.pickup_datetime.dt.day,
                    month = df.pickup_datetime.dt.month,
                    year = df.pickup_datetime.dt.year,
                    dayofweek = df.pickup_datetime.dt.dayofweek)

```

```

[15]: df = df.drop(["pickup_datetime"], axis =1)
      df

```

```

[15]:      fare_amount  pickup_longitude  pickup_latitude  dropoff_longitude \
0           7.5      -73.999817      40.738354      -73.999512
1           7.7      -73.994355      40.728225      -73.994710
2          12.9      -74.005043      40.740770      -73.962565
3           5.3      -73.976124      40.790844      -73.965316
4          16.0      -73.925023      40.744085      -73.973082
...          ...          ...          ...          ...

```

199995	3.0	-73.987042	40.739367	-73.986525
199996	7.5	-73.984722	40.736837	-74.006672
199997	30.9	-73.986017	40.756487	-73.858957
199998	14.5	-73.997124	40.725452	-73.983215
199999	14.1	-73.984395	40.720077	-73.985508

	dropoff_latitude	passenger_count	hour	day	month	year	dayofweek
0	40.723217	1	19	7	5	2015	3
1	40.750325	1	20	17	7	2009	4
2	40.772647	1	21	24	8	2009	0
3	40.803349	3	8	26	6	2009	4
4	40.761247	5	17	28	8	2014	3
...
199995	40.740297	1	10	28	10	2012	6
199996	40.739620	1	1	14	3	2014	4
199997	40.692588	2	0	29	6	2009	0
199998	40.695415	1	14	20	5	2015	2
199999	40.768793	1	4	15	5	2010	5

[200000 rows x 11 columns]

```
[16]: from math import *

def distance_formula(longitude1, latitude1, longitude2, latitude2):
    travel_dist = []

    for pos in range(len(longitude1)):
        lon1, lan1, lon2, lan2 = map(radians, [longitude1[pos], latitude1[pos],
        ↪ longitude2[pos], latitude2[pos]])
        dist_lon = lon2 - lon1
        dist_lan = lan2 - lan1

        a = sin(dist_lan/2)**2 + cos(lan1) * cos(lan2) * sin(dist_lon/2)**2

        c = 2 * asin(sqrt(a)) * 6371
        travel_dist.append(c)

    return travel_dist
```

```
[17]: df['dist_travel_km'] = distance_formula(df.pickup_longitude.to_numpy(), df.
    ↪ pickup_latitude.to_numpy(), df.dropoff_longitude.to_numpy(), df.
    ↪ dropoff_latitude.to_numpy())
```

```
[18]: def remove_outlier(df1 , col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
    IQR = Q3 - Q1
```

```

lower = Q1-1.5*IQR
upper= Q3+1.5*IQR
df[col] = np.clip(df1[col] , lower , upper)
return df1

def remove_all(df1 , col_list):
    for i in col_list:
        df1 = remove_outlier(df , i)
    return df1

```

```
[19]: df = remove_all(df , df.iloc[:, 0::])
```

```
[20]: df_x =
    ↪df[['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude', 'passenger_co
df_y = df['fare_amount']
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
scaler.fit_transform(df_x)
```

```
[20]: array([[ -1.04868902, -0.46151906, -1.01920818, ...,  1.75477984,
           -0.02487235, -0.52650782],
           [-0.78236166, -0.85626623, -0.80934206, ..., -1.4772954 ,
            0.48875385, -0.17833351],
           [-1.30351906, -0.3673497 ,  0.59560201, ..., -1.4772954 ,
           -1.56575094,  0.98130178],
           ...,
           [-0.37579257,  0.24518925,  2.36704659, ..., -1.4772954 ,
           -1.56575094,  2.25341213],
           [-0.91736709, -0.96432192, -0.30695085, ...,  1.75477984,
           -0.53849854,  0.30827932],
           [-0.29670226, -1.17381839, -0.40715524, ..., -0.9386162 ,
            1.00238004,  1.15281348]])
```

```
[21]: x_train, x_test, y_train, y_test = train_test_split(df_x, df_y, test_size=0.2,
    ↪random_state=1)
```

```
[22]: from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score, mean_squared_error
```

```
[23]: reg = LinearRegression()
reg.fit(x_train, y_train)
```

```
[23]: LinearRegression()
```

```
[24]: y_pred_lin = reg.predict(x_test)
print(y_pred_lin)
```

```
[ 6.2761843  5.09988121  9.43640959 ... 11.07661434 12.15390374
 11.41498106]
```

```
[25]: rmse_lr = np.sqrt(mean_squared_error(y_test, y_pred_lin))
      r2_lr = r2_score(y_test, y_pred_lin)

      print("Linear Regression - RMSE:", rmse_lr)
      print("Linear Regression - R2 score:", r2_lr)
```

```
Linear Regression - RMSE: 2.7039561758284734
Linear Regression - R2 score: 0.75390636301046
```

```
[26]: from sklearn.ensemble import RandomForestRegressor
```

```
[27]: rf = RandomForestRegressor(n_estimators=100)
      rf.fit(x_train, y_train)
```

```
[27]: RandomForestRegressor()
```

```
[28]: RandomForestRegressor()
```

```
[28]: RandomForestRegressor()
```

```
[29]: y_pred_rf = rf.predict(x_test)
      print(y_pred_rf)
```

```
[ 4.999  6.674  9.2175 ... 11.975  11.012  13.429 ]
```

```
[30]: rmse_rf= np.sqrt(mean_squared_error(y_test, y_pred_rf))
      r2_rf = r2_score(y_test, y_pred_rf)

      print("Random Forest - RMSE:", rmse_rf)
      print("Random Forest - R2 score:", r2_rf)
```

```
Random Forest - RMSE: 2.364547091246532
Random Forest - R2 score: 0.8118097917070926
```

```
[ ]:
```